

```
In [148]: import pandas as pd
```

```
In [149]: import warnings  
warnings.filterwarnings('ignore')
```

```
In [150]: data=pd.read_csv("/home/placement/Downloads/fiat500.csv")
```

```
In [151]: data.describe()
```

```
Out[151]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [152]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column             Non-Null Count  Dtype  
---  -
0   ID                  1538 non-null  int64  
1   model               1538 non-null  object  
2   engine_power        1538 non-null  int64  
3   age_in_days         1538 non-null  int64  
4   km                  1538 non-null  int64  
5   previous_owners     1538 non-null  int64  
6   lat                 1538 non-null  float64 
7   lon                 1538 non-null  float64 
8   price               1538 non-null  int64  
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [153]: data1=data.drop(['ID','lat','lon'],axis=1)

```
In [154]: data1
```

```
Out[154]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [155]: data2=data.loc[(data.model=='lounge')]
```

In [156]: data2

Out[156]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
6	7	lounge	51	731	11600	1	44.907242	8.611560	10750
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9190
11	12	lounge	51	366	17500	1	45.069679	7.704920	10990
...
1528	1529	lounge	51	2861	126000	1	43.841980	10.515310	5500
1529	1530	lounge	51	731	22551	1	38.122070	13.361120	9900
1530	1531	lounge	51	670	29000	1	45.764648	8.994500	10800
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990

1094 rows × 9 columns

In [157]: data1=pd.get_dummies(data)

In [158]: data1

Out[158]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price	model_lounge	model_pop	model_sport
0	1	51	882	25000	1	44.907242	8.611560	8900	1	0	0
1	2	51	1186	32500	1	45.666359	12.241890	8800	0	1	0
2	3	74	4658	142228	1	45.503300	11.417840	4200	0	0	1
3	4	51	2739	160000	1	40.633171	17.634609	6000	1	0	0
4	5	73	3074	106880	1	41.903221	12.495650	5700	0	1	0
...
1533	1534	51	3712	115280	1	45.069679	7.704920	5200	0	0	1
1534	1535	74	3835	112000	1	45.845692	8.666870	4600	1	0	0
1535	1536	51	2223	60457	1	45.481541	9.413480	7500	0	1	0
1536	1537	51	2557	80750	1	45.000702	7.682270	5990	1	0	0
1537	1538	51	1766	54276	1	40.323410	17.568270	7900	0	1	0

1538 rows × 11 columns

In [159]: data1.shape

Out[159]: (1538, 11)

In [160]: y=data1['price']
x=data1.drop('price',axis=1)

In [161]:

y

```
Out[161]: 0      8900
          1      8800
          2      4200
          3      6000
          4      5700
          ...
          1533    5200
          1534    4600
          1535    7500
          1536    5990
          1537    7900
```

Name: price, Length: 1538, dtype: int64

In [162]:

x

Out[162]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	model_lounge	model_pop	model_sport
0	1	51	882	25000	1	44.907242	8.611560	1	0	0
1	2	51	1186	32500	1	45.666359	12.241890	0	1	0
2	3	74	4658	142228	1	45.503300	11.417840	0	0	1
3	4	51	2739	160000	1	40.633171	17.634609	1	0	0
4	5	73	3074	106880	1	41.903221	12.495650	0	1	0
...
1533	1534	51	3712	115280	1	45.069679	7.704920	0	0	1
1534	1535	74	3835	112000	1	45.845692	8.666870	1	0	0
1535	1536	51	2223	60457	1	45.481541	9.413480	0	1	0
1536	1537	51	2557	80750	1	45.000702	7.682270	1	0	0
1537	1538	51	1766	54276	1	40.323410	17.568270	0	1	0

1538 rows × 10 columns

```
In [163]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [164]: x_test.head()
```

```
Out[164]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	model_lounge	model_pop	model_sport
481	482	51	3197	120000	2	40.174702	18.167629	0	1	0
76	77	62	2101	103000	1	45.797859	8.644440	0	1	0
1502	1503	51	670	32473	1	41.107880	14.208810	1	0	0
669	670	51	913	29000	1	45.778591	8.946250	1	0	0
1409	1410	51	762	18800	1	45.538689	9.928310	1	0	0

```
In [165]: y_test.head()
```

```
Out[165]: 481    7900
76    7900
1502   9400
669    8500
1409   9700
Name: price, dtype: int64
```

```
In [166]: x_train.head()
```

```
Out[166]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	model_lounge	model_pop	model_sport
527	528	51	425	13111	1	45.022388	7.58602	1	0	0
129	130	51	1127	21400	1	44.332531	7.54592	1	0	0
602	603	51	2039	57039	1	40.748241	14.52835	0	1	0
331	332	51	1155	40700	1	42.143860	12.54016	1	0	0
323	324	51	425	16783	1	41.903221	12.49565	1	0	0

```
In [167]: y_train.head()
```

```
Out[167]: 527    9990
          129    9500
          602    7590
          331    8750
          323    9100
          Name: price, dtype: int64
```

```
In [169]: from sklearn.linear_model import LinearRegression
          reg=LinearRegression()
          reg.fit(x_train,y_train)
```

```
Out[169]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [170]: ypred=reg.predict(x_test)
```


In [171]: ypred

```
Out[171]: array([[ 5819.19308764,  7248.82914161,  9741.8936974 ,  9798.98033074,
        10055.00624601,  9551.4955679 ,  9758.01743879, 10122.9778365 ,
        9654.9661814 ,  9251.1403257 , 10478.09512253,  7807.3005255 ,
        7705.15873781,  6295.63244894,  9545.40486313, 10422.92177704,
        9616.90811615,  7756.9171161 ,  4893.88454414, 10581.46142719,
        10465.24078346, 10443.29318231,  7518.43696046, 10028.21911459,
        6990.73118896,  8989.86900819,  4823.51364349,  6989.03118684,
        7822.83203734,  9683.17944083,  7344.21343132,  5341.43860798,
        5420.78405336,  5092.38401339,  8971.44357515,  5702.81242412,
        9920.16285466,  8334.58448277,  6220.93323723,  8389.23958511,
        9695.84208061,  6859.59630725,  9101.22635456, 10063.22592995,
        8621.83915759, 10175.06753933,  9063.21918346,  8867.24865352,
        7094.44228184,  9058.37693565,  9474.82390731, 10406.09102832,
        10112.65006224,  6820.90463865,  9700.36507783,  9382.18149429,
        9632.57617775, 10553.81356008,  9847.21129432,  7247.16814789,
        9990.23331336,  7084.23300123,  9977.34233656,  7245.01115798,
        6490.89305576,  9737.86785115,  9853.54349825,  8568.7125607 ,
        8506.81438703,  6484.69051659,  7883.1895563 ,  6870.28308427,
        8263.36833348, 10551.03496347,  7434.71134313,  8637.85174602,
        8762.87817027, 10010.47000077,  7324.60000000,  8527.72426022])
```

In [172]: `from sklearn.metrics import r2_score`
`r2_score (y_test,ypred)`

Out[172]: 0.8428319728488683

In [173]: `from sklearn.metrics import mean_squared_error`
`mean_squared_error(ypred,y_test)`

Out[173]: 577189.6736608233

In [174]: `import math`
`a=577189.6736608233`
`print(math.sqrt(a))`

759.7300005007195

```
In [175]: Results=pd.DataFrame(columns=['price', 'predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(15)
```

Out[175]:

	index	price	predicted	Id
0	481	7900	5819.193088	0
1	76	7900	7248.829142	1
2	1502	9400	9741.893697	2
3	669	8500	9798.980331	3
4	1409	9700	10055.006246	4
5	1414	9900	9551.495568	5
6	1089	9900	9758.017439	6
7	1507	9950	10122.977837	7
8	970	10700	9654.966181	8
9	1198	8999	9251.140326	9
10	1088	9890	10478.095123	10
11	576	7990	7807.300526	11
12	965	7380	7705.158738	12
13	1488	6800	6295.632449	13
14	1432	8900	9545.404863	14

```
In [176]: from sklearn.model_selection import GridSearchCV
          from sklearn.linear_model import Ridge

          alpha = [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20, 30]

          ridge = Ridge()

          parameters = {'alpha': alpha}

          ridge_regressor = GridSearchCV(ridge, parameters)

          ridge_regressor.fit(x_train, y_train)
```

```
Out[176]: GridSearchCV(estimator=Ridge(),
                        param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                              5, 10, 20, 30]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [177]: ridge_regressor.best_params_
```

```
Out[177]: {'alpha': 30}
```

```
In [178]: ypred=ridge_regressor.predict(x_test)
```

In [179]: ypred

```
Out[179]: array([[ 5819.29853963,  7264.57491791,  9738.88270579,  9794.47839507,
        10050.35072397,  9548.82126342,  9750.20283681, 10118.76944676,
        9656.2363147 ,  9247.20526971, 10474.46954   ,  7800.22901654,
        7695.77817109,  6317.72759712,  9542.31591964, 10420.56401819,
        9641.41151189,  7746.00639631,  4909.8836403 , 10575.9590411 ,
        10447.28811339, 10440.93292264,  7492.38838409, 10023.77691534,
        6993.68728961,  9012.08291864,  4822.19326199,  6980.28023648,
        7816.98010837,  9677.40953383,  7336.83285116,  5328.48898357,
        5430.56724837,  5073.55505956,  8936.61797371,  5693.27871402,
        9939.05168141,  8309.31737942,  6213.4861627 ,  8409.23567191,
        9693.0168448 ,  6873.88260945,  9125.50501592, 10086.14831587,
        8615.86224143, 10176.53613302,  9084.51866352,  8863.55539547,
        7082.99237491,  9054.89062122,  9472.31391242, 10401.71056408,
        10110.24411314,  6837.93965878,  9697.35103307,  9405.38998926,
        9653.58056603, 10549.50117936,  9840.96037592,  7263.00085407,
        9985.59074909,  7075.08499779,  9973.28948478,  7235.14597592,
        6479.76739925,  9738.88533685,  9849.24462578,  8587.72121603,
        8500.85578807,  6473.43468594,  7872.26497993,  6871.81868177,
        8258.28361541, 10545.55397382,  7426.43055818,  8629.44613409,
        8750.28575402, 10007.64052610,  7310.04778021,  8524.18104124])
```

```
In [180]: ridge=Ridge(alpha=30)
          ridge.fit(x_train,y_train)
          y_pred_ridge=ridge.predict(x_test)
```

```
In [181]: from sklearn.metrics import mean_squared_error
          Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
          Ridge_Error
```

```
Out[181]: 574728.5696156605
```

```
In [182]: from sklearn.metrics import r2_score
          r2_score (y_test,y_pred_ridge)
```

```
Out[182]: 0.8435021284061197
```

```
In [184]: Results=pd.DataFrame(columns=['Actual','predicted'])
Results['Actual']=y_test
Results['Predicted']=ypred
Results=Results.reset_index()
Results['Id']=Results.index
Results.head(15)
```

Out[184]:

	index	Actual	predicted	Predicted	Id
0	481	7900	NaN	5819.298540	0
1	76	7900	NaN	7264.574918	1
2	1502	9400	NaN	9738.882706	2
3	669	8500	NaN	9794.478395	3
4	1409	9700	NaN	10050.350724	4
5	1414	9900	NaN	9548.821263	5
6	1089	9900	NaN	9750.202837	6
7	1507	9950	NaN	10118.769447	7
8	970	10700	NaN	9656.236315	8
9	1198	8999	NaN	9247.205270	9
10	1088	9890	NaN	10474.469540	10
11	576	7990	NaN	7800.229017	11
12	965	7380	NaN	7695.778171	12
13	1488	6800	NaN	6317.727597	13
14	1432	8900	NaN	9542.315920	14

```
In [185]: from sklearn.linear_model import ElasticNet
          from sklearn.model_selection import GridSearchCV

          elastic = ElasticNet()

          parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-4, 1e-3, 1e-2, 1, 5, 10, 20]}

          elastic_regressor = GridSearchCV(elastic, parameters)

          elastic_regressor.fit(x_train, y_train)
```

```
Out[185]: GridSearchCV(estimator=ElasticNet(),
                        param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 1,
                                                5, 10, 20]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [186]: elastic_regressor.best_params_
```

```
Out[186]: {'alpha': 0.01}
```

```
In [188]: elastic=ElasticNet(alpha=.01)
          elastic.fit(x_train,y_train)
          y_pred_elastic=elastic.predict(x_test)
```

```
In [189]: y_pred_elastic
```

```
Out[189]: array([ 5818.88155606, 7251.60740134, 9741.37548309, 9798.17563202,  
 10054.18656906, 9551.03709384, 9756.60308305, 10122.24386003,  
 9655.23079132, 9250.43675403, 10477.46946422, 7805.98706383,  
 7703.41076005, 6299.65982145, 9544.86952947, 10422.52128897,  
 9622.08375661, 7754.80253699, 4897.38480628, 10580.47637966,  
 10461.71330862, 10442.89200323, 7512.81872938, 10027.42619282,  
 6990.9795985 , 8993.93966292, 4822.9274055 , 6987.38971162,  
 7821.74361059, 9682.14401582, 7342.84571741, 5338.98378321,  
 5422.36411921, 5088.69094616, 8964.54155191, 5701.00133273,  
 9923.63970799, 8329.65576797, 6219.5421374 , 8392.88958456,  
 9695.35125166, 6862.17433129, 9106.3590491 , 10067.46187307,  
 8620.73186465, 10175.36762819, 9067.11601592, 8866.57789002,  
 7092.28977184, 9057.74520111, 9474.38711846, 10405.31440756,  
 10112.23689137, 6823.98653466, 9699.84010068, 9386.45543089,  
 9637.09234337, 10553.05108845, 9846.08937468, 7250.04609281,  
 9989.41187642, 7082.51579887, 9976.61490792, 7243.17422389,  
 6488.79520193, 9738.09464099, 9852.77779753, 8572.18811531,  
 8505.72502159, 6482.56240442, 7881.15444662, 6870.90829793,  
 8262.43238047, 10550.05609925, 7433.16232868, 8636.30752919,  
 8761.50105130, 10000.07027067, 7322.62650502, 8520.65020602])
```

```
In [190]: from sklearn.metrics import r2_score  
r2_score(y_test,y_pred_elastic)
```

```
Out[190]: 0.8429739684420192
```

```
In [191]: from sklearn.metrics import mean_squared_error  
elastic_Error=mean_squared_error(y_pred_elastic,y_test)  
elastic_Error
```

```
Out[191]: 576668.2037947337
```

```
In [192]: Results=pd.DataFrame(columns=['Actual','predicted'])
Results['Actual']=y_test
Results['Predicted']=y_pred_elastic
Results=Results.reset_index()
Results['Id']=Results.index
Results
```

```
Out[192]:
```

	index	Actual	predicted	Predicted	Id
0	481	7900	NaN	5818.881556	0
1	76	7900	NaN	7251.607401	1
2	1502	9400	NaN	9741.375483	2
3	669	8500	NaN	9798.175632	3
4	1409	9700	NaN	10054.186569	4
...
503	291	10900	NaN	10120.713199	503
504	596	5699	NaN	6291.601668	504
505	1489	9500	NaN	10020.222896	505
506	1436	6990	NaN	8247.810365	506
507	575	10900	NaN	10337.015702	507

508 rows × 5 columns

```
In [ ]:
```